



SUPPORT VECTOR MACHINES

E. Fersini

INTRODUCTION

- During this lab, we will illustrate the following:
 1. How to **train** a SVM
 2. Observing how the **choice of cost** can affect SVM
 3. Visualizing **SVM fit**
 4. **Predicting the labels** of a testing dataset based on the model trained by SVM
 5. **Tuning SVM**
- The two most well known and popular support vector machine tools are **libsvm** in the e1071 package and **SVMlite** in the klaR package.
 - During this lab we will use **libsvm**
 - We will work on the **iris dataset**

TRAINING SVM

- Load the e1071 package:

```
> library(e1071)
```

- Load the dataset, create training and testing:

```
> ind = sample(2, nrow(iris), replace = TRUE, prob=c(0.7, 0.3))  
> testset = iris[ind == 2,]  
> trainset = iris[ind == 1,]
```

- Train the model:


```
> svm.model = svm(Species ~ ., data=trainset, kernel='linear', cost=1)
```



TRAINING SVM

- Our SVM constructs a hyperplane (or set of hyperplanes) that maximize the margin width between two classes.

$$\min_{w,b} \frac{1}{2} \|w\|^2 + C \sum_i \xi_i$$



TRAINING SVM

- Our SVM constructs a hyperplane (or set of hyperplanes) that maximize the margin width between two classes.
- We can obtain overall information about the induced model with the `print()` function:

```
> print(svm.model)
```

- The SVM model has a **cost parameter C** that controls the trade-off between the slack variable penalty (misclassifications) and width of the margin:
 - a **small cost** creates a large margin (a soft margin) and **allows more misclassifications**;
 - a **large cost** creates a narrow margin (a hard margin) and **allows few misclassifications**.
 - $C=\text{infinity}$ implies that all the constraints must be satisfied: no points within the margin

CHOOSE the COST of SVM

- With the following steps, we will understand how large and small cost will affect SVM.
- Performing the following steps to generate two different classification examples with different costs:
 1. Subset the iris dataset with columns named as Sepal.Length, Sepal.Width, Species, **only** with species in setosa and virginica

```
> iris.subset = subset(trainset, select=c("Sepal.Length", "Sepal.Width",  
"Species"), Species %in% c("setosa","virginica"))
```

CHOOSE the COST of SVM

2. Then, you can generate a scatter plot with Sepal.Length as the x-axis and the Sepal.Width as the y-axis:

```
> plot(x=iris.subset$Sepal.Length,y=iris.subset$Sepal.Width,  
col=iris.subset$Species, pch=19)
```

3. Next, we can train SVM based on iris.subset with the **cost** equal to 1:

```
> svm.model = svm(Species ~ ., data=iris.subset, kernel='linear', cost=1,  
scale=FALSE)
```

CHOOSE the COST of SVM

4. Then, we can mark the support vectors with blue circles:

```
> points(iris.subset[svm.model$index,c(1,2)],col="blue",cex=2)
```

5. We can add a separation line on the plot:

```
> w = t(svm.model$coefs) %*% svm.model$SV  
> c = -svm.model$rho  
> abline (a=-c/w[1,2],b=-w[1,1]/w[1,2], col="red", lty=5)
```

intercept

slope

dotted line

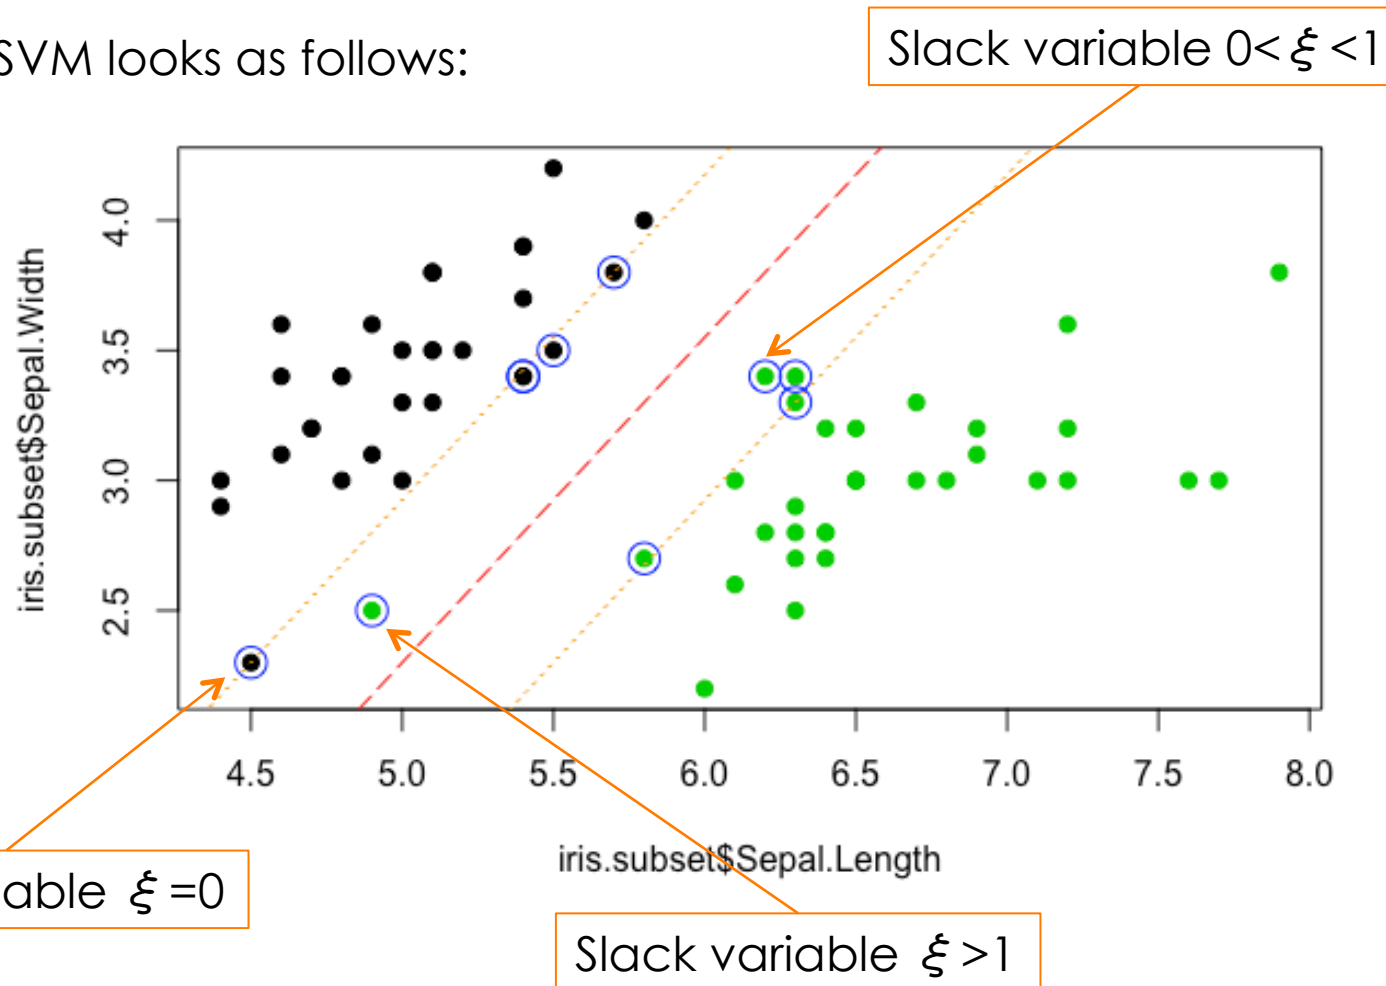
CHOOSE the COST of SVM

5. Lastly, we can add the margin on the plot:

```
> abline(a=(-c-1)/w[1,2], b=-w[1,1]/w[1,2], col="orange", lty=3)  
> abline(a=(-c+1)/w[1,2], b=-w[1,1]/w[1,2], col="orange", lty=3)
```

CHOOSE the COST of SVM

- Our SVM looks as follows:



CHOOSE the COST of SVM

- We can now work with a different cost parameter $c=10,000$
 1. Reset the plot:

```
> plot(x=iris.subset$Sepal.Length,y=iris.subset$Sepal.Width,  
col=iris.subset$Species, pch=19)
```

2. Train SVM with $c=10000$:

```
> svm.model = svm(Species ~ ., data=iris.subset, kernel='linear',  
cost=10000, scale=FALSE)
```

CHOOSE the COST of SVM

3. Mark the support vectors with blue circles:

```
> points(iris.subset[svm.model$index,c(1,2)],col="blue",cex=2)
```


4. Add a separation line on the plot:

```
> w = t(svm.model$coefs) %*% svm.model$SV  
> c = -svm.model$rho  
> abline (a=-c/w[1,2],b=-w[1,1]/w[1,2], col="red", lty=5)
```

```
> abline(a=(-c-1)/w[1,2], b=-w[1,1]/w[1,2], col="orange", lty=3)  
> abline(a=(-c+1)/w[1,2], b=-w[1,1]/w[1,2], col="orange", lty=3)
```

CHOOSE the COST of SVM

- Now, our SVM with $c=10,000$ looks as follows:



PREDICTION

- We can predict the label of new instances (testset) by using the already trained SVM model:

```
> svm.pred = predict(svm.model, testset)
```

- And create our confusion matrix:

```
> svm.table=table(svm.pred, testset$Species)  
> svm.table
```

ASSIGNMENT

- Investigate the parameters:
 - **kernel**: radial, linear, polynomial, radial basis, sigmoid.
 - **cost**: the default value is 1; the larger the value, the smaller the margin is.
 - **gamma**: the default value is equal to $(1/\text{data dimension})$; increasing gamma implies an increasing number of support vectors.
 - it controls the shape of the separating hyperplane

ASSIGNMENT

- Define an SVM suitable for the “iris” dataset with the optimal parameter setting
 - Perform automatic tuning
 - With all the classes


```
> tuned = tune.svm(Species ~ ., data = iris.subset, kernel='linear',  
cost=c(0.001, 0.01, 0.1, 1, 5, 10, 100))
```


DOING MULTI-CLASS CLASSIFICATION

- SVMs can only handle two-class outputs (i.e. a categorical output variable with arity 2).
- How to handle multiple classes
 - E.g., classify documents into three categories: *sports, business, politics*
- Solutions:
 - **one-vs-all**, learn N SVM's
 - SVM 1 learns "Output==1" vs "Output != 1"
 - SVM 2 learns "Output==2" vs "Output != 2"
 - ...
 - SVM N learns "Output==N" vs "Output != N"
 - **Pairwise coupling**, learn N SVM's
 - SVM 1 learns "Output==1" vs "Output 2"
 - SVM 2 learns "Output==1" vs "Output != 3"
 - ...
 - SVM N learns "Output==2" vs "Output != 3"
 -


ONE-VS-ALL

-  vs the other classes: $\text{red}(d)$




ONE-VS-ALL

- ■ vs the other classes: red(d)
- ■ vs the other classes: blue (d)



ONE-VS-ALL

- ■ vs the other classes: $\text{red}(d)$
- ■ vs the other classes: $\text{blue}(d)$
- ■ vs the other classes: $\text{cyan}(d)$



ONE-VS-ALL

- vs the other classes: red(d)
- vs the other classes: blue(d)
- vs the other classes: cyan(d)


Given a test sample d' , how to decide its color ?

*Assign d to the color function with the **largest score!***

$$\mathbf{w}^T \mathbf{x}_i + b$$


PAIRWISE COUPLING

-  VS 









PAIRWISE COUPLING

- ■ VS ■
• ■ VS ■ -----



PAIRWISE COUPLING

-  VS 
-  VS  - - - - -
-  VS  - . - . - .
- ...

Which class for  ?

Assign to the class that wins
the most pairwise comparison

